

# INTRO TO CONCURRENCY

Created by [Bahram Aghaei](#)

## **COMMON I/O OPERATIONS IN WEB APPLICATIONS:**

- Downloading the content of a web page
- Communicating over a network
- Running several queries against a database

But, the I/O operations tend to be **slow** for different reasons.

# WHAT WE'LL COVER

- What is `asyncio`?
- Difference between CPU-bound vs. I/O-bound
- Concurrency, Parallelism, and Multitasking
- Global Interpreter Lock (GIL)
- Non-blocking I/O + Event Loop

# WHAT IS ASYNCIO?

In a synchronous app, code runs sequentially, waiting for each task to finish.

**Concurrency** allows multiple tasks to run simultaneously, keeping the app responsive.

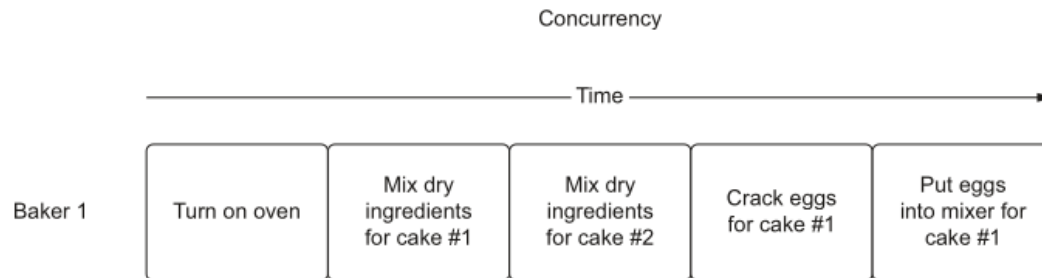
Asynchronous programming lets tasks run in the background, freeing up the system to do other work.

`asyncio` is a Python library for managing tasks asynchronously using a **single-threaded event loop**.

**CPU-BOUND**  
**VS**  
**I/O-BOUND**

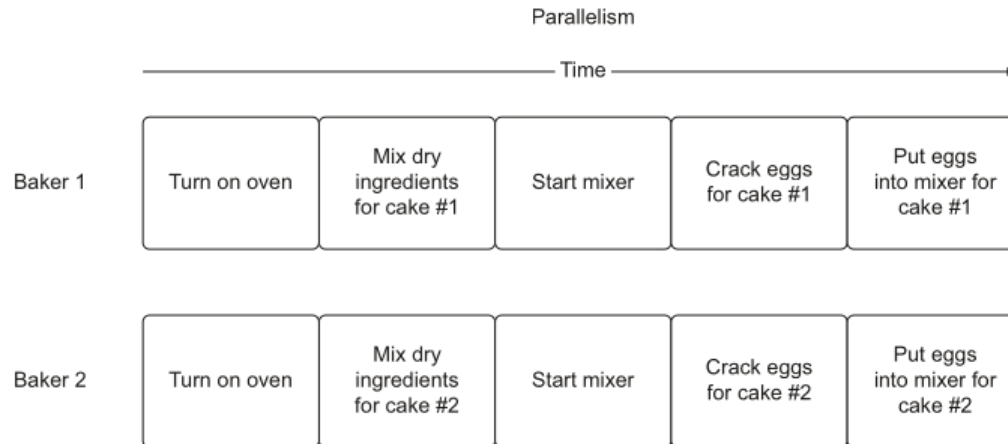
# CONCURRENCY

When we say two tasks are happening concurrently, we mean those tasks are happening at the same time.



# PARALLELISM

When we say something is running in **parallel**, we mean not only are there **two or more tasks** happening concurrently, but they are also **executing at the same time**.





# MULTITASKING

- Preemptive Multitasking
- Cooperative Multitasking

# PREEMPTIVE MULTITASKING

In this model, the operating system decides how to switch between tasks using time slicing, known as **preempting**.

# COOPERATIVE MULTITASKING

In this model, the program itself defines points where other tasks can run.

**“I’m pausing my task; go ahead and run other tasks.”**

- Less resource intensive
- More granularity

# GLOBAL INTERPRETER LOCK (GIL)

GIL prevents one Python process from **executing more than one Python bytecode instruction** at any given time.

## Why does the GIL exist?

- due to how memory is managed in CPython.
- CPython is not thread safe.

**Takeaway:** For CPU-bound tasks, multithreading is not beneficial.

# GIL (CONTINUED)

**Is the GIL ever released?**

- The GIL is released when I/O operations occur.

**Why release the GIL for I/O but not CPU-bound tasks?**

- I/O operations use system calls outside of the Python runtime, allowing the GIL to be released.
- CPU-bound tasks do not benefit from GIL release because they run within Python.

# GIL (CONTINUED)

## Using asyncio with GIL:

- asyncio uses **coroutines**, which are lightweight threads.
- **Important:** asyncio does not circumvent the GIL; it operates within its constraints.

# NON-BLOCKING I/O + EVENT LOOP

Sockets are blocking by default.

```
1 import socket
2
3 sock = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
4
5 # Connect the socket to the server's address and port
6 server_address = ('localhost', 8080)
7 sock.connect(server_address)
8
9 # Blocking read from the socket
10 data = sock.recv(1024) # Blocking call
11 print(f"Received: {data.decode()}")
```

## SOLUTION 🤔

At the OS level, we can operate in non-blocking mode using **event notification systems**.

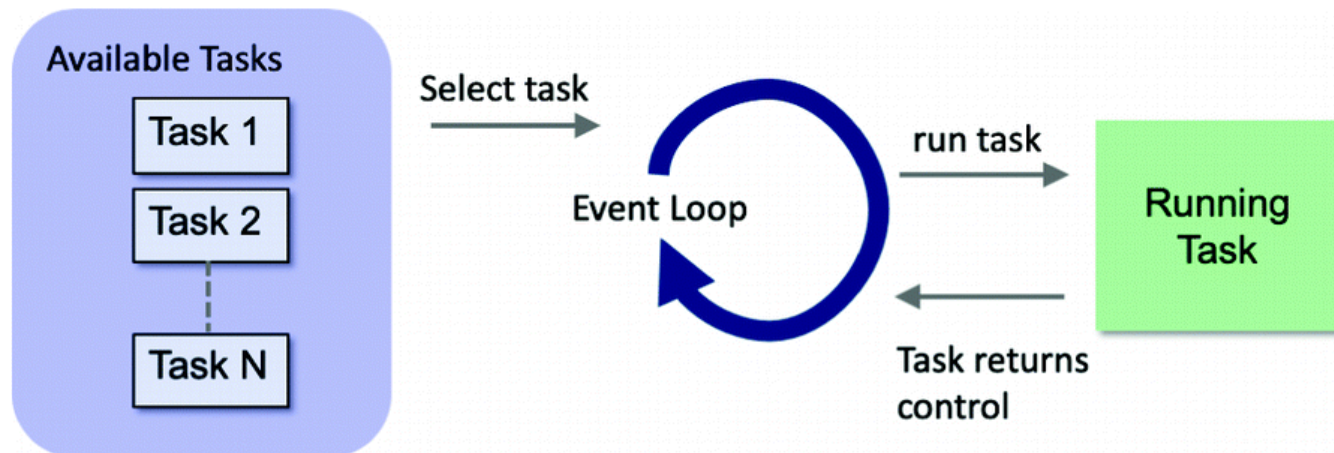
- **kqueue** — FreeBSD and MacOS
- **epoll** — Linux
- **IOCP (I/O completion port)** — Windows



# THE EVENT LOOP: MANAGING I/O TASKS

How do we track which tasks are waiting for I/O  
versus those that are regular Python code?

**an event loop**



# BASIC EVENT LOOP EXAMPLE

```
1 from collections import deque
2 messages = deque()
3
4 while True:
5     if messages:
6         message = messages.pop()
7         process_message(message)
8
9 def make_request():
10     cpu_bound_setup()
11     io_bound_web_request()
12     cpu_bound_postprocess()
13
14 task_one = make_request()
15 task_two = make_request()
```

